

# CPSC 483

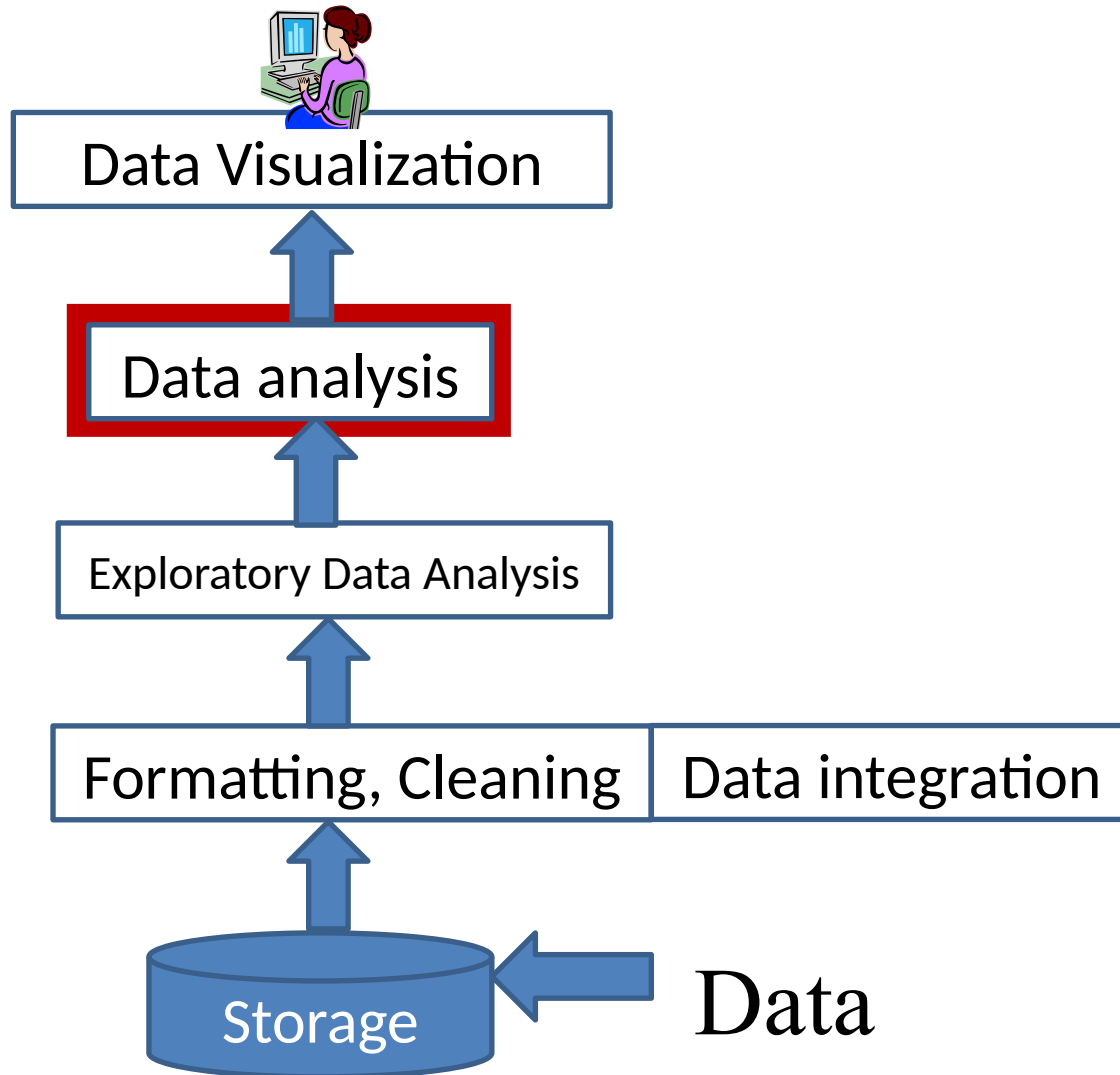
## Evaluation of ML models

Anand Panangadan  
apanangadan@fullerton.edu

# What we will cover this week

- Linear regression
- Evaluating a regression model
- Analytical and iterative approaches
- Regularization

# The Data Science Process




# What is a Math Model?

- Model
  - Representation of a phenomenon
  - Describes the relationship between variables
- Mathematical model
  - Numerically describe relationship between variables

“All models are wrong, but some are useful”

# Regression Model

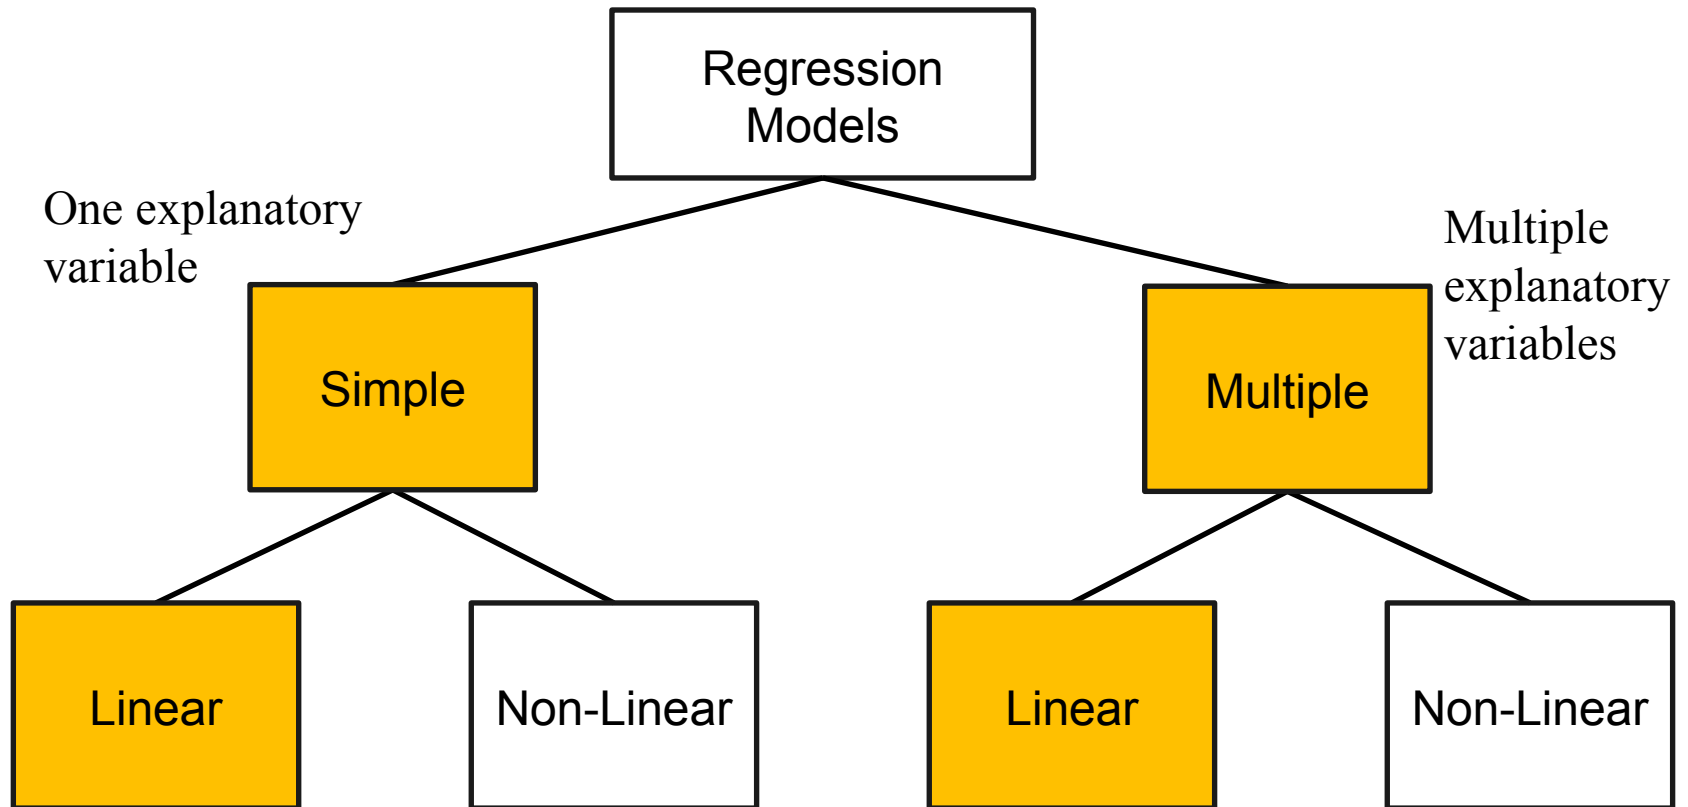
- Relationship between one **dependent variable** and **explanatory variable(s)**
- Dependent variable is **continuous**
- One equation describes the relationship



# Steps to modeling

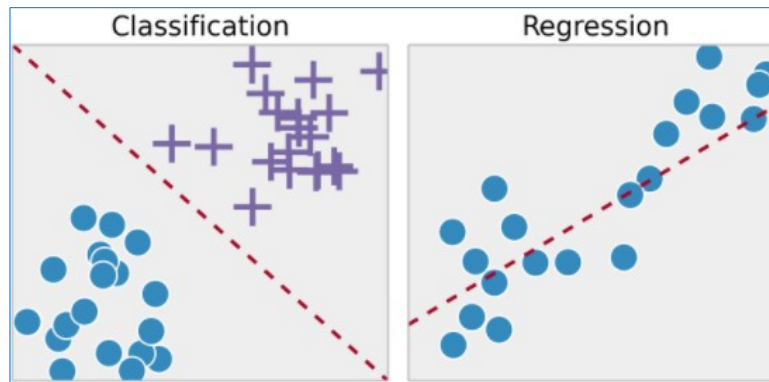
1. Define a family of models
  - Specify the generic pattern of relationship between variables
    - Linear, Quadratic, ...
2. Fit the model to the data
  - Identify the “best” model from the family of models
  - Give values to the model parameters

# Types of Regression Models



# Supervised Learning

- Given a training data set
  - Learn a **function**
    - The process of learning is called **training**.
  - Predict outcomes for a given (new data):
    - If the learned **function**  $f$  is used to predict a continuous value, it's called **regression**.
    - If the learned **classifier**  $f$  is used to determine a discrete value, it's called **classification**.
- **Classification vs. regression**
  - **Classification** is to find the **decision boundary** that **separates** the different groups of data as clearly as possible.
  - **Regression** is to find the **function** that **fits** the data the best.






# Class work

- Consider the problem of estimating Height given Weight and Age
  - What are some linear models for this?
  - What are some nonlinear models for this?

# Linear Regression Model

- Relationship between variables is a **linear function**



# Interpretation of Coefficients

## 1. Slope

- per unit change in the dependent variable for each unit change in the independent variable.


## 2. Y-Intercept

- Average Value of  $Y$  When  $X = 0$

# Class work

- How does Petal.Length change with Petal.Width in the iris dataset?
    - Which is the dependent variable? Which is the explanatory/independent variable?
    - Draw a scatterplot
    - Draw *any* straight line that fits the data well
- ```
from sklearn.datasets import load_iris
iris = load_iris(as_frame=True)
fig, ax = plt.subplots()
...
plt.plot(x, y, "b.")
plt.grid()
ax.axline(xy1=(0, intercept), slope=slope, color='red')
```

# Which model/line is better?




# Least Squares (LS)

- Best Fit
  - Select the slope & intercept that gives the smallest difference between Actual Y values () & Predicted Y values ().
  - Square the differences
  - Differences are called *errors* or *residuals*

$$\sum_{i=1}^n \left( Y_i - \hat{Y}_i \right)^2 = \sum_{i=1}^n \epsilon_i^2$$

# Least Squares Graphically

LS minimizes



# LS minimization is efficient!

- Sample slope
- Sample  $Y$  – intercept
- Prediction equation



# When **can** we use LM?

- One dependent variable
- Dependent variable is continuous

# Linear Regression in scikit-learn

- `from sklearn.linear_model import LinearRegression`
- `lin_reg = LinearRegression()`
- `lin_reg.fit(X, y)`
- `lin_reg.intercept_, lin_reg.coef_`

# Plot the model

```
y_predict = lin_reg.predict(X)  
plt.plot(X, y_predict, "r-")
```

OR for simple linear regression:

```
ax.axline(xy1=(0, lin_reg.intercept_),  
slope=lin_reg.coef_[0], color='red')
```

# Class work

- Consider the “Auto MPG” dataset which “concerns city-cycle fuel consumption in miles per gallon, to be predicted in terms of 3 multivalued discrete and 5 continuous attributes.” The goal is to model mpg given engine displacement.
  - Load the **autompg.csv** file on Canvas
  - Which is the dependent variable? Which is the independent variable?
  - Plot mpg vs. displacement (code, plot)
  - Overlay best fit line over the dataset (code, plot)


# When **should** we use LM?

- Is the “best” fit model a **good** model?!
- Evaluating the model
- Two conflicting objectives
  - Goodness-of-Fit
    - We want model to match the data
  - Complexity
    - We want model to be “simple”
- “Principle of parsimony”
  - Find a model that is as simple as possible without sacrificing too much goodness-of-fit

# Coefficient of Determination

- The proportion of total variation (SST) that is explained by the regression (SSR) is the **Coefficient of Determination** ( $R^2$ )
- ranges between 0 and 1
  - higher its value, the more accurate is the regression model

# Sum of Squares of Error (SSE)




**A least squares regression selects the line with the lowest total sum of squared prediction errors.**

**This value is called the Sum of Squares of Error (SSE)**

**This is the “**unexplained**” variation**

# Sum of Squares Regression (SSR)



**The Sum of Squares Regression (SSR) is the sum of the squared differences between the prediction for each observation and the population mean.**

**This is the “**explained**” variation**



# Total Sum of Squares (SST)

**SSR = (measure of explained variation)**

**SSE = (measure of unexplained variation)**

**SST = SSR + SSE = (measure of total variation in y)**

# Coefficient of Determination

```
lin_reg.score(X, y)
```

OR

```
from sklearn.metrics import r2_score  
r2_score(y, y_predict)
```

# Regression Diagnostics

- The three conditions required for the validity of the regression analysis are:
  1. the error variable is normally distributed
  2. the error variance is independent of  $x$
  3. The errors are independent of each other
- How can we diagnose violations of these conditions?

# Residuals

- Also called “errors”
- $\text{residuals} = y - y_{\text{predict}}$

# Residual Analysis


- Examine the residuals
- Non-normality.
  - Examine the residual histogram and look for a bell shaped curve with a mean close to zero

```
residuals = y - y_predict  
plt.hist(residuals) # optional bins= parameter
```

# Residual Analysis

- Plot a scatterplot
  - Residuals vs  $x$
  - Points should appear to be randomly scattered around zero


```
plt.plot(X, residuals, 'b.')
```




Residuals appear randomly scattered around zero, and their spread appears constant



*homoscedasticity*



systematic behavior in the residuals



Residuals scattered randomly about zero but variability isn't constant



*heteroscedasticity*

# Outliers

- Outlier: an observation that is unusually small or large
- Several possibilities need to be investigated when an outlier is observed:
  - There was an error in recording the value.
  - The point does not belong in the sample.
  - The observation is valid.
- *Suspect* an observation is an outlier if its  
 $| \text{standardized residual} | > 2$



# Making predictions

- The fitted model can be used to make predictions for new  $x$  values
  - For accuracy,  $x$  should be within the range seen in the data used for modeling (interpolation)

# Procedure for Regression Diagnostics

- Gather data for the variables in the model
- Draw a scatterplot to determine whether a linear model appears to be appropriate
- Determine the regression equation
- Check the required conditions for the errors (residual analysis)
- Check the existence of outliers
- Assess the model fit ( $R^2$ )
- If the model fits the data well, use the model to predict new values

# Class work

- Load the **autompg.csv** file on Canvas
- The goal is to model mpg given engine displacement
- Predict the mpg of a car with engine displacement=250
- Predict the mpg of a car with engine displacement=600

# Why is LR important?

- Simple
- Efficient
- Assumptions are reasonable
- Model is surprisingly powerful
  - Multiple predictor variables
  - Can use transformed predictors

# Multiple Linear Regression

- More than one independent variable can be used to explain variance in the dependent variable
- Multiple regression takes the form
- $k+1$ : number of model parameters to estimate

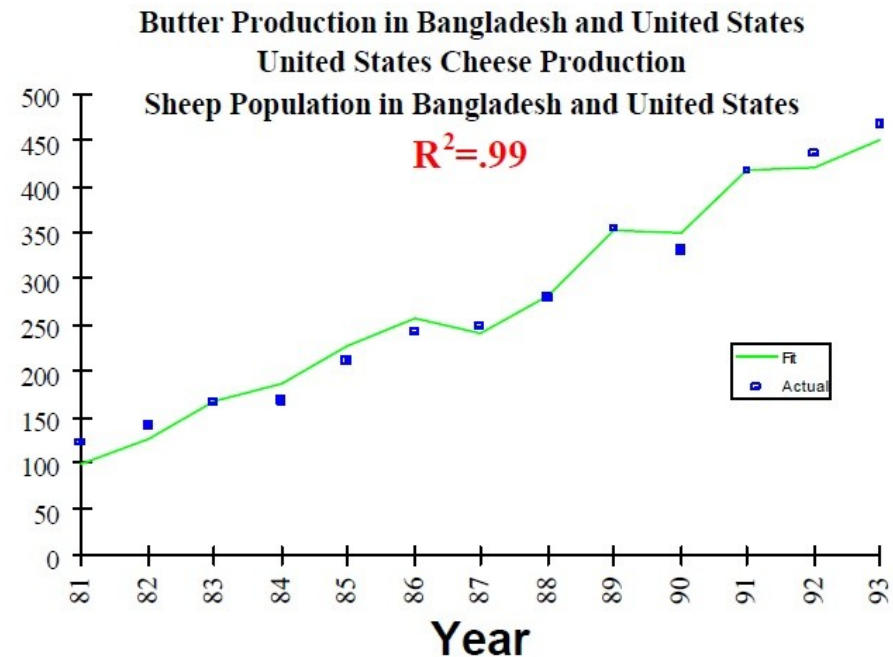
# Classwork

- Write the best fit linear model equation relating Petal.Width to Petal.Length and Sepal.Length in the iris dataset
  - Write code to get coefficients

# Adjusted R2

- Will adding more independent variables *always* increase R2?
- What if the independent variables are completely random?
- Surprisingly, yes!

- Must focus on number of variables
  - More variables ☹ less trustworthy model
- Adjusted R2
  - Adjusts for the number of variables by reducing R2 for more complex models
- Overfitting: A particular problem in Big Data
- [How to get a high R2 model for the S&P 500?!](#)



# Adjusted R<sup>2</sup>

$$\bar{R}^2 = 1 - (1 - R^2) \frac{n - 1}{n - p - 1}$$

- $p$  = total number of explanatory variables (excluding the intercept)
- $n$  = sample size



# The Linear Regression Model

- The **general linear model** can be written as an equation:
  - are independent (exploratory) variables/features, representing the data.
  - is a dependent (response) variable.
  - are coefficients/weights;  $k$  is number of features.
  - $\epsilon$  is a random variable (noise by measurements or others), assumed to have a  $N(0, \sigma^2)$  distribution;  $\mu$  is a random variable with mean,  $\mu$  and variance  $\sigma^2$ .

| <i>Price (K\$) (Y)</i> | <i>Size (X<sub>1</sub>)</i> | <i>Bedroom (X<sub>2</sub>)</i> | <i>Bathroom (X<sub>3</sub>)</i> | <i>Built year (X<sub>4</sub>)</i> |
|------------------------|-----------------------------|--------------------------------|---------------------------------|-----------------------------------|
| 375 ()                 | 1024 (x <sub>11</sub> )     | 3 (x <sub>12</sub> )           | 2 (x <sub>13</sub> )            | 1978 (x <sub>14</sub> )           |
| 425 ()                 | 1329 (x <sub>21</sub> )     | 3 (x <sub>22</sub> )           | 5 (x <sub>23</sub> )            | 1992 (x <sub>24</sub> )           |
| ...                    | ...                         | ...                            | ...                             | ...                               |
| 465 ()                 | 1893 (x <sub>N1</sub> )     | 4 (x <sub>N2</sub> )           | 4 (x <sub>N3</sub> )            | 1980 (x <sub>N4</sub> )           |

Uppercase is a column data.

Lowercase is a row data.

is a value.

|  |
|--|
|  |
|  |
|  |
|  |

ing dataset.

# Regression Model in Vector/Matrix Form

- **Stacking all the equations**

- k number of variables (features)
- N number of equations (rows)

|  |
|--|
|  |
|  |
|  |
|  |

- **The equations in vector form**


- or

—

- **The equations in matrix form**

- => or

is a feature (column) vector.  
is a row vector (ith training example).



# The Goal of Regression Modeling

- **Housing price data set**

| <i>Price (K\$) (<math>Y</math>)</i> | <i>Size (<math>X_1</math>)</i> | <i>Bedroom (<math>X_2</math>)</i> | <i>Bathroom (<math>X_3</math>)</i> | <i>Built year (<math>X_4</math>)</i> |
|-------------------------------------|--------------------------------|-----------------------------------|------------------------------------|--------------------------------------|
| 375 ( $y_1$ )                       | 1024 ( $x_{11}$ )              | 3 ( $x_{12}$ )                    | 2 ( $x_{13}$ )                     | 1978 ( $x_{14}$ )                    |
| 425 ( $y_2$ )                       | 1329 ( $x_{21}$ )              | 3 ( $x_{22}$ )                    | 5 ( $x_{23}$ )                     | 1992 ( $x_{24}$ )                    |
| ...                                 | ...                            | ...                               | ...                                | ...                                  |
| 465 ( $y_N$ )                       | 1893 ( $x_{N1}$ )              | 4 ( $x_{N2}$ )                    | 4 ( $x_{N3}$ )                     | 1980 ( $x_{N4}$ )                    |

- **Housing price data set in matrix form**

- ,  $X$  and  $Y$  are from training data (known);  $w$  are unknown.

- **The goal is to estimate** of the function

# FINDING REGRESSION MODELS ANALYTICALLY

# The Least Squares Approach


- The **basic idea** of the (ordinary) least squares approach

- Find  $w$  that minimizes the squared error:

is the actual value.

is the estimated parameter.

is the predicted value of  $y$ , for a given data  $x_i$ .



- Finding  $w$  with minimum squared distances between data in a training set and predicted line.

- **Note:** The regression problem is **restated** as an **optimization problem** of an **error function**.

- **Why least square?**

- This problem can be solved **analytically** (convincing probabilistic interpretation).  
How?


# Solving Linear Regression by the Least Square Approach

- Find parameters  $w$  that minimizes the loss function for  $\langle x, y \rangle$  value pairs of all the training examples (Gauss and Legendre, 1809).

$$\text{Minimize } l(w) = \frac{1}{N} \sum_i (y_i - f_w(x_i))^2$$

**Squaring** residuals is important!

$N$  is the number of samples in a training data set.





– Residual sum of Squares (RSS) =

$$\text{Loss} = \text{Mean Squared Error (MSE)} = \frac{1}{N} \sum_i (y_i - w^T x_i)^2$$

$\text{Argmin}_w l(w)$

# Derivatives to Test Turning Points



- The **first derivative** test:
  - A **local minimum** is where  $f'(x)$  changes from **-** to **+**.
  - A **local maximum** is where  $f'(x)$  changes from **+** to **-**.
  - **Solving**  $f'(x) = 0$  returns the x-coordinates of **all turning points** (solutions).
- The **second derivative** test:
  - If  $f''(x) > 0$ , the turning point at x is concave up (**convex**);  $f(x)$  has a local **minimum** at x.
  - **Else if**  $f''(x) < 0$ , the turning point at x is concave down (concave);  $f(x)$  has a local **maximum** at x.
  - **Else:** Saddle point (neither local minima nor maxima)

# Deriving a Multiple Regression Model

$$\text{Argmin}_w l(w) = \frac{1}{N} \sum_i (y_i - w^T x_i)^2 \quad \text{from}$$

- To find  $w$  that minimizes
  - Re-arranging  $\Rightarrow$   
 $\Rightarrow$
  - Taking a **partial derivative** of w.r.t  $w$  and solving it ( $= 0$ )  
  
(Refer to Rogers et al. book for the detailed derivation)
  - To get the **estimated  $w$** , multiply both sides:
  - Finally, we derived a linear model:
    - We analytically found a solution that satisfies **the least square error**.



# Predicting a New Value

- Given a **new example** , the **predicted value**, from the **model** obtained by **the least square** is computed by:

where

- A value of  $\hat{y}$  is computed by **dot product** of  $\mathbf{x}$  and  $\mathbf{w}$ , that is,  $\hat{y} = \mathbf{x} \cdot \mathbf{w}$

$$\hat{y} = x_{i1} + x_{i2} + \dots + x_{ik}$$

# FINDING REGRESSION MODELS ITERATIVELY

# Optimizing Loss Function Iteratively

(to deal with a **huge matrix X** from a large data set and computing **inverse matrix**)

- For **simple regression**, **minimize** w.r.t

$$l(w_0, w_1) = \frac{1}{N} \sum_i (y_i - f_w(x_i))^2 \text{ or } \frac{1}{N} \sum_i (y_i - (w_0 + x_i * w_1))^2$$

- **Iterative method** for a **computational solution**

- i. **Start** with some **initial** parameters e.g.,
- ii. **Keep changing** to **reduce**
- iii. **Stop when** desirable conditions are satisfied.


- **Questions:**

- 1) What can be **initial values** of parameters ?
- 2) How to **change** the parameter values , by **how much**?
- 3) How do we know if the **new parameters w** will **reduce** ?
- 4) What are the conditions to **stop** the **iteration**?

# (Steepest) Gradient Descent Method

- To **minimize** we can use the following iterative method
  - Start with an **initial**
  - Keep changing  $w_0$  to **reduce** :  $w_{\text{next}} = w_0 \pm \Delta w$
  - Stop when** certain conditions are satisfied.

How to tell whether the **gradient** at a point  $w_0$  is negative or positive?





To **reduce** , left case:  $\Delta w$  should be positive, but right case:  $\Delta w$  negative.

How do we determine the **sign of  $\Delta w$**  (e.g., how do we change , subtract or add)?

# Linear Regression as a Convex Optimization Problem

**Gradient descent** (also called *steepest gradient descent*) is an **iterative optimization approach** to find the **minimum** or **maximum** of a function by taking **1<sup>st</sup> order derivative**.



**Gradient descent** will reach a **global optimum** if a function has only one optimum. Otherwise, it will reach a **local optimum**.

# Finding $w$ by Gradient Descent Method

To **minimize** for simple regression, **minimize** and .

## To find

- Initialize  $w$  with a random value
- **Repeat** updating **until** convergence

where hyper parameter is a **tuning parameter** (or learning rate) to control the **speed** and **accuracy** of convergence.

## To find

- Initialize  $w$  with a random value
- **Repeat** updating **until** convergence

## Derivatives of $J$ to determine $\Delta w$

- **From Loss Function**

Or instead of MSE, we can use  $RSS =$

- ▶ The **partial derivatives** of  $J$  w.r.t.  $w$  and  $b$ :

$=$

$=$

# Simple Regression by Gradient Descent Method

**Initialize** with random values

**Repeat until** convergence {

Predicted value

Actual value

Each  $(w_0, w_1)$  is calculated by the residual sum of all the N examples in a training data set.

}

- At each iteration, **update**  $\Delta w$  and **simultaneously** (independently).

- A tuning parameter is set empirically (e.g., [0.0, 1.0]) as it can inflate or deflate  $\Delta w$ , impacting convergence speed and possibly accuracy.
  - Too small, slow convergence
  - Too large, fail to converge or diverge



# Multiple Regression by Gradient Descent Method

$$Y = X_0 + X_1 + X_2 + X_3 + \dots + ( = 1)$$

**Initialize** with k+1 random values (k number of features)

**Repeat** until convergence {

for **each** (j=0, ..., k) in parallel

Update each **simultaneously** (independently)

}

For example, we can compute  $w_0, w_1, w_2$ :

– (= 1)

Each is calculated by the residual sum of all the N examples in a training data set.

# A Variant of Gradient Descent Method

- **Batch Gradient Descent** (we just discussed this method.)
  - At each step, use **ALL the training samples**.

$$w_j = w_j - \alpha \frac{1}{N} \left[ \sum_i (w^T x_i - y_i) x_{ij} \right]$$

- **Stochastic Gradient Descent**
  - Use a **portion of random training samples** at each iteration.
  - This is **effective** with **big data** or for **online/adaptive learning**.

# When to Update the Weights

- **Batch training**
  - All the training examples are used in every iteration. The weight update for each epoch for the average error. Faster convergence to a local minimum since the weights are moved in the direction that most of the inputs want them to move.
- **Stochastic Gradient Descent (SGD)**
  - The weights are updated for each one randomly selected input. Repeat the process until convergence. Often used when the training set is very large.
- **Minibatches Gradient Descent (MGD)**
  - Combination of batch and SGD
  - The training set are split into fixed-size or random batches, then update the weights for each batch.
  - The training set are then randomly shuffled into new batches and the process is repeated.
  - If the batches are small, then the global minimum may be found although at the cost of heading in the wrong direction.

# The Least Square vs. Gradient Descent


- **The least square method** (analytical solution for minimum error)
  - Very simple formula and simple implementation
  - **but** for big data, a large matrix requires large memory
  - **but** computing matrix inverse can be expensive for a large matrix
- **Gradient descent method** (iterative solution by approximation)
  - Fast even with big data (because computing is dot product of vectors)
  - *Stochastic gradient descent* is very *memory efficient*.
  - The method is easily extensible to other problems.
  - **but** it requires parameter tweaking:
    - How to decide convergence?
    - What is the best learning rate?
- **Most mathematical approaches for machine learning problems**
  - solved based on **analytical solution** (very few), **estimation**, **approximation**, and (min/max) **optimization**.

# Class work

- Load the **autompg.csv** file on Canvas
- The goal is to model mpg given engine displacement
- Use Stochastic gradient descent to predict
  - the mpg of a car with engine displacement=250
  - the mpg of a car with engine displacement=600
- Experiment with different values of hyperparameters
  - max\_iter, eta0
- Are your prediction results sensitive to the value of these hyperparameters?
- How do these predictions compare to those obtained analytically?

```
from sklearn.linear_model import SGDRegressor  
sgd_reg = SGDRegressor(max_iter=1000, tol=1e-5, penalty=None,  
eta0=0.01)
```

# Transforming predictors




**Nonlinear functions can also be fit as regressions. Common choices include Power, Logarithmic, Exponential, and Logistic, but any continuous function can be used.**

# Transforming predictors

- Create a new independent variable from existing variables
  - Eg.:  $z = x^2$
- Set up new regression equation
  - $y =$
  - Equivalent to:
- Can even include multiple variables
  - Eg.:

# Transforming predictors





# Class work: nonlinear transformation

For the iris dataset

- Build a linear model to predict Petal.Length given:
  - $\text{Petal.Width}^2$
- What is the regression equation?
- Plot the model predictions

# Classwork: nonlinear transformation

For the autmpg dataset (on Canvas)

- Build a linear model to predict mpg given:
  - $1/(\text{displacement}^2)$
- What is the regression equation?
- Plot the model predictions

# Transforming predictors: categorical variables

- Transform categorical variables using one-hot encoding **dummy variables**
- If categorical variable takes  $n$  values, create
  - $n - 1$  0/1 variables

# Transforming predictors: categorical variables

- Example: Species variable in iris dataset
- Factor with 3 levels: *setosa*, *versicolor*, *virginica*
- First dummy variable for setosa
  - =1 when Species=setosa
  - =0 otherwise
- Second dummy variable for versicolor
  - =1 when Species=versicolor
  - =0 otherwise
- Third dummy variable for virginica
  - =1 when Species= virginica
  - =0 otherwise

| Value      |   |   |   |
|------------|---|---|---|
| Setosa     | 1 | 0 | 0 |
| Versicolor | 0 | 1 | 0 |
| Virginica  | 0 | 0 | 1 |

# Multiple parallel lines with a categorical variable

- Build a linear model to predict Petal.Length given Petal.Width and Species
- What is the linear regression equation?

# Multiple parallel lines with a categorical variable

- Build a linear model to predict Petal.Length given Petal.Width and Species
- What is the linear regression equation?
- $PL = 1.211397 + 1.018712 * \text{Petal.Width} + 1.697791 * \text{Speciesversicolor} + 2.276693 * \text{Speciesvirginica}$

IF Species == setosa (both dummy vars = 0,0)

- $PL = 1.211397 + 1.018712 * \text{Petal.Width}$

IF Species == versicolor (dummy vars = 1,0)

- $PL = (1.211397 + 1.697791) + 1.018712 * \text{Petal.Width}$
- $PL = 2.909188 + 1.018712 * \text{Petal.Width}$

IF Species == virginica (dummy vars = 0,1)

- $PL = (1.211397 + 2.276693) + 1.018712 * \text{Petal Width}$

# Multiple parallel lines with a categorical variable

- Build a linear model to predict Petal.Length given Petal.Width and Species
- What is the linear regression equation?
- $\text{Petal.Length} = 1.21140 + (1.01871 * \text{Petal.Width}) + (1.69779 * \text{Speciesversicolor}) + (2.27669 * \text{Speciesvirginica})$
- If (Species == setosa)
  - $\text{Petal.Length} = 1.21140 + (1.01871 * \text{Petal.Width}) + \cancel{(1.69779 * \text{Speciesversicolor})} + \cancel{(2.27669 * \text{Speciesvirginica})}$
  - $\text{Petal.Length} = 1.21140 + (1.01871 * \text{Petal.Width})$
- If (Species == versicolor)
  - $\text{Petal.Length} = 1.21140 + (1.01871 * \text{Petal.Width}) + (1.69779 * 1) + \cancel{(2.27669 * \text{Speciesvirginica})}$
  - $\text{Petal.Length} = (1.21140 + 1.69779) + (1.01871 * \text{Petal.Width})$
- If (Species == virginica)
  - $\text{Petal.Length} = 1.21140 + (1.01871 * \text{Petal.Width}) + \cancel{(1.69779 * 0)} + (2.27669 * 1)$
  - $\text{Petal.Length} = (1.21140 + 2.27669) + (1.01871 * \text{Petal.Width})$

# Visualizing the linear model with categorical variables

- Multiple lines
  - Each dummy variable creates its own line
- Example: Species variable in iris dataset
- Plot the predictions



# Classwork


Consider the **autompg.csv** dataset on Canvas. The goal is to model mpg given engine displacement (a continuous variable) and number of cylinders (pretend this is a categorical variable).

- Which is the dependent variable? Which are the independent variables?
- Load the **autompg.csv** file on Canvas and convert cylinders variable to a factor
- Create a linear model called `mod_displ_cyl` of mpg vs. displacement and cylinders.
- Give the model equations relating mpg with displacement and cylinders.
- Plot mpg vs. displacement and overlay the best fit model.


# REGULARIZING COMPLEX MODELS TO AVOID OVERFITTING

# Choosing the Best Model


(a) A linear model with  $L = 1.358$



(b) An 8<sup>th</sup> order polynomial model with  $L = 0.459$



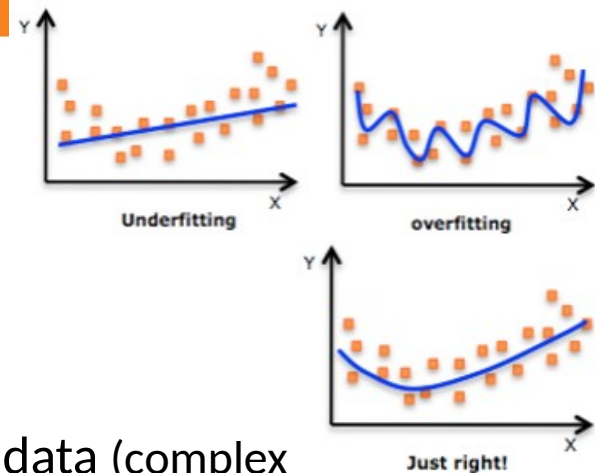
(c) A customized model with  $L = 1.1037$



- Which model looks the best?

- The models with increasing complexity (**higher order model**) fit better, resulting in lower RSS values as it gets closer to the training data, but **some predictions can be wildly inaccurate** for **even small deviations** from the training data instances.

# Over-fitting and Under-fitting






- Each model is a **generalization** from the **training data**
- **Over-fitting**: A model fits too closely to the training data (complex but the quality of the predictions can be poor).
  - **High variance**: **Low training error** and **high generalization error**, causing an algorithm to model the random noise or memorize the training data.
  - **Occam's Razor**: Given two models with the same generalization errors, the simpler model is preferred over the more complex model.
- **Under-fitting**: A model cannot adequately capture the underlying structure of the data.
  - **High bias**: **High training error** and **high systemic error**, causing an algorithm to miss the relevant relations between the independent and dependent variables.

# Bias-Variance trade-off

- A model's generalization error can be expressed as the sum of three types of errors:
- Bias
  - Error due to wrong assumptions, such as assuming that the data is linear when it is actually quadratic.
  - A high-bias model is most likely to **underfit** the training data.
- Variance
  - Model's sensitivity to variations in the training data.
  - A model with many degrees of freedom (e.g., high-degree polynomial) is likely to have high variance and thus **overfit** the training data.
- Irreducible error
  - Noise in the data. Nothing that the model can do anything about.

# Bias-Variance trade-off

- Increasing a model's complexity will typically increase its variance and reduce its bias.
- Conversely, reducing a model's complexity increases its bias and reduces its variance.
- Tradeoff is to determine the optimal model complexity such that it is able to generalize well without over-fitting.
- This is a challenging task!




# Underfitting and Overfitting Problems


- **Common causes: recommended solutions**
  - Data quality and random noise: preprocess the data
  - Lack of representative examples: use more training data
    - How much data are appropriate? A rule of thumb for statistically reliable result: At least 5 training examples for each dimension and >30 examples
    - Generate more training data using semi-supervised learning or generative method
  - Too many variables: reduce the dimension and select the core features
  - Algorithms with poor generalization capability: use a better algorithm
- **Common causes of underfitting: recommended solutions**
  - Too simple models: make the models complex
- **Common causes of overfitting: recommended solutions**
  - Too complex models: simplify the models (regularization), early stopping, dropout, pruning, Bayesian priors, etc. and evaluate its performance on unseen data (testing/validation data).

# Complex Models


(a) A linear model with  $l = 1.358$



(b) An 8<sup>th</sup> order polynomial model with  $l = 0.459$



(c) A customized model with  $l = 1.1037$



- What's going on with higher order models?
  - The models with increasing complexity (**higher order model**) **fit better** the data and result in **lower RSS values** (*low error in training*) and but possibly matching *even the noise in the data*. Therefore, *some predictions may be very inaccurate* with **small deviations** in the new examples (**not properly generalized**).
  - To correct this problem, add a **penalty term** in the **loss function** (**regularization**).



# Regularizing Models

- **Regularizing** a model:

$M(w) + \lambda R(w)$  where  $\lambda R(w)$  is a regularized term (penalty term or shrinkage factor);  
 $\lambda$  can be any real value  $>0$ ;  $w$  are shrinking coefficients.

- Why penalize the magnitude of coefficients  $w$ ?

- **Regularizing** (note: this is a **loss function**)

- LASSO (Least Absolute Shrinkage and Selection Operator):

$(\lambda ||w_i||_1, \text{L1 regularization})$


- Ridge:  $(\lambda ||w_i||_2, \text{L2 regularization})$

- Elastic net: LASSO + Ridge

- **The role of the tuning parameter  $\lambda$**

- If  $\lambda = 0$ , linear regression (without regularization).
- The larger  $\lambda$ , more aggressively penalization is, the coefficient is close to zero, resulting in a simpler model.
- $\lambda$  should be chosen wisely (using cross-validation) as **it controls** the **tradeoff** between **penalizing not fitting** the data and **penalizing overly complex models**.

# The Effects of the Tuning Parameter $\lambda$




The **higher** the **sum** in **w**, the **more complex** the model. To penalize the complex model, keep the **w** value lower by increasing  **$\lambda$** .

Varying the regularization parameter  **$\lambda$**  for a 5<sup>th</sup> order polynomial function

- If  **$\lambda = 0$** , we get the original solution.
- If  **$\lambda$  is too small** (1e-06), the model is likely to be too complex, getting closer to the general shape of the exact 5<sup>th</sup> order polynomial function.
- If  **$\lambda$  is too large**, less complex but may not capture any useful structure of the data.
- We need to determine the best value of  **$\lambda$** .

# Ridge regression




Linear (left) and a polynomial (right) models with various levels of ridge regularization

# Ridge Regression


- **Ridge regression using the least square method**
  - , solving for  $w$ :  $= 0$
  - From , the regularized least square solution is:  
(analytical solution)
- **Ridge regression using the gradient descent method**
  - Adding the regularization term to
  - $+ ]$
  - $+$

# LASSO Regression



- LASSO regression:
  - No derivative at  $x = 0$ , we use “coordinate descent” that optimizes each parameter separately, holding all the others fixed (based on the concept of subgradients).
  - Unlike Ridge, no analytical solution since the solution is nonlinear in  $y$ .
- Derivation of LASSO regression by Coordinate Descent
  - $RSS = \Rightarrow$
  - Let  $\beta$ , then since (if normalized).
  - The regularized term,  $-\lambda$  when  $\beta < 0$ ,  $[-\lambda, \lambda]$  when  $\beta = 0$ , and  $\lambda$  when  $\beta > 0$  (subgradients)
  - For minima of  $\beta$ , let  $\beta = 0$ , then solve for :
    - Case  $\beta < 0$ :  $\beta = 0$ ,  $\beta = +$ . Need  $\beta < 0$
    - Case  $\beta = 0$ :  $[-2 - \lambda, -2 + \lambda]$  to contain 0. Need  $\lambda \leq 2$
    - Case  $\beta > 0$ :  $\beta = 0$ ,  $\beta = -$ . Need  $\beta > 0$



# LASSO regression



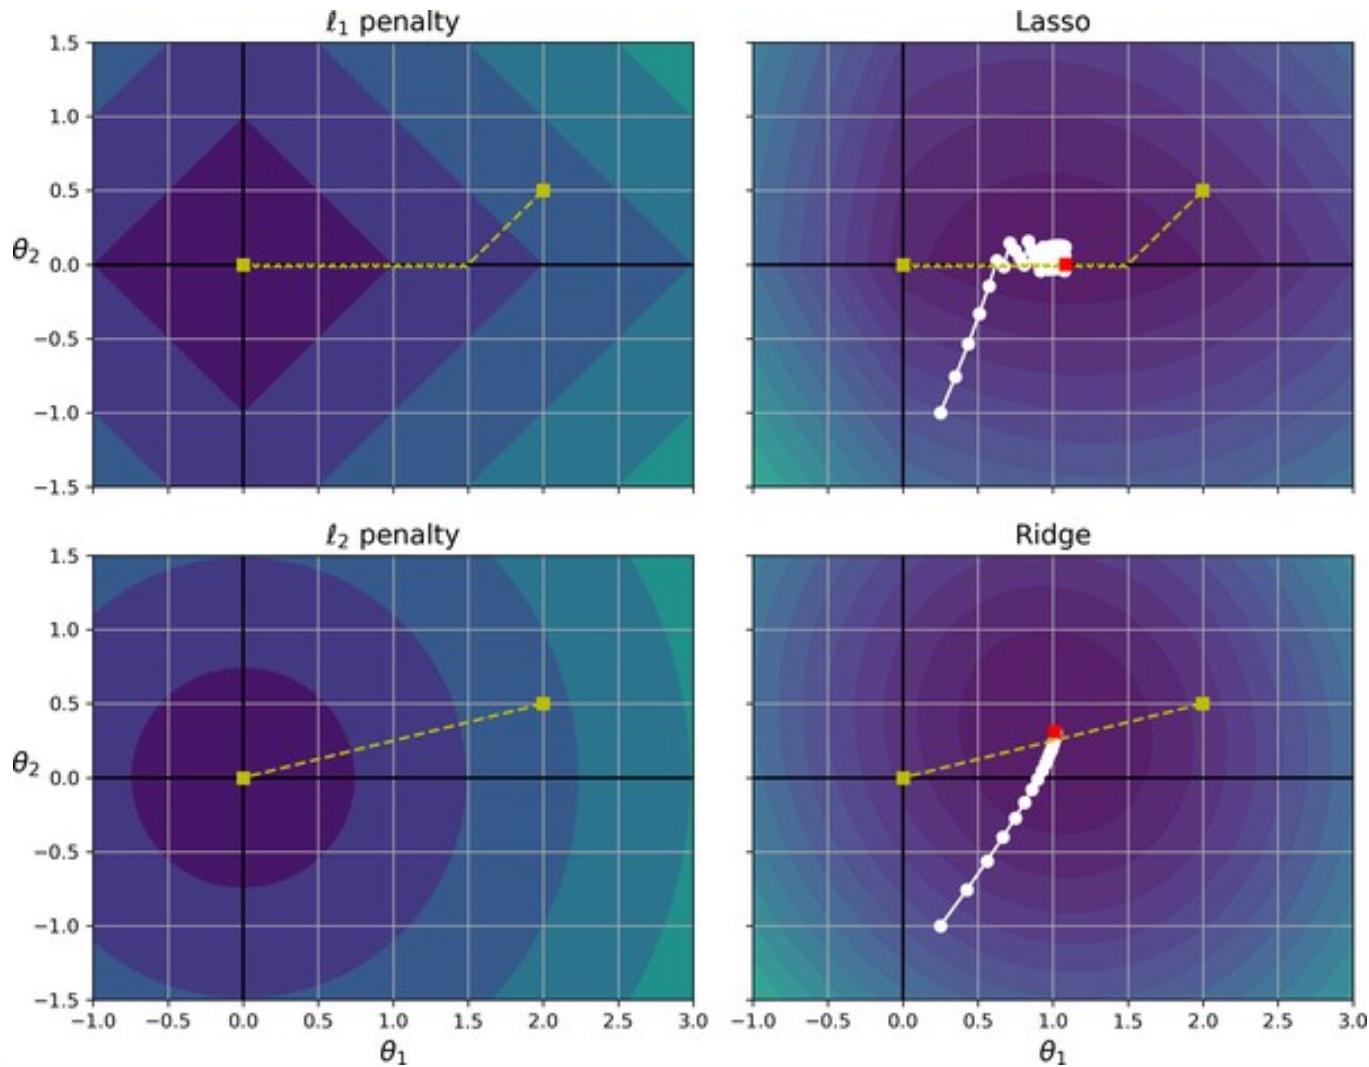
Linear (left) and a polynomial (right) models with various levels of LASSO regularization

# Ridge vs LASSO

- LASSO and Ridge regressions



# Ridge vs LASSO






# Ridge vs LASSO

- **The effects of  $\lambda$** 
  - As  $\lambda$  increases, **Ridge** can shrink the coefficients asymptotically close to zero, while LASSO can shrink them all the way to zero, eliminating the less important features.
    - If the squared sum in LASSO hits one of the corners, the coefficients are shrunk to zero.
  - For the same  $\lambda$ , **LASSO** has much smaller coefficients and higher RSS compared to Ridge. Many coefficients are zero even for very small  $\lambda$ , “sparsity”.
  - One way to find a proper  $\lambda$ , start with a relatively large  $\lambda$ , then decrease it slowly.
  - The coefficients of correlated predictors are similar in Ridge while one for the correlated predictors has a larger coefficient and the rest are nearly zeroed in LASSO.
  - LASSO works well for a model with high multicollinearity and can be useful for feature selection, but it may be too dependent on data and thus unstable.
- Overall neither one is better. How about LASSO + Ridge?


# Elastic Net Regressio



- **Elastic net** (for both feature selection and regularization)

- For  $\alpha = [0,1]$ ,  $\lambda \geq 0$ ,

,



- When  $\alpha = 1$ , Elastic net  $\equiv$  LASSO.
  - As  $\alpha$  shrinks toward 0, it approaches Ridge.
  - For other values of  $\alpha$ , the penalty term interpolates between L1 and the L2 of w.

- **Update rule**

- Find the coefficients of Ridge, then perform the LASSO on the Ridge regression coefficients to shrink the coefficients.

# Class work

- Load the iris dataset
- The goal is to model Petal.Width given Petal.Length, Sepal.Length, and Sepal.Width
- Use LASSO regression and experiment with different values of the hyperparameter alpha
- What are the coefficients?
- What is the linear regression equation?
- Is your equation sensitive to the value of alpha?

```
from sklearn.linear_model import Lasso
lasso_reg = Lasso(alpha=0.1)
lasso_reg.fit(X, y)
```

# Logistic Regression


- Modify linear regression to do classification
- Model estimates the probability that an instance belongs to a class.
  - If the estimated probability  $> 0.5$ , then predict that the instance belongs to that class (positive)
  - Otherwise ( $p < 0.5$ ), predict that it does not belong (negative)
- Binary classifier

# Logistic Regression

- Model computes
  - a weighted sum of the input features
  - Followed by the **logistic** of this weighted sum
- **Logistic function:**
  - A sigmoid function that outputs a number between 0 and 1 for any input

$$\hat{p} = h_{\theta}(\mathbf{x}) = \sigma(\theta^T \mathbf{x})$$

$$\sigma(t) = \frac{1}{1 + \exp(-t)}$$



# Class work

- Load the iris dataset
- The goal is to classify instances as Species="virginica" or not given Petal.Length, Sepal.Length, Sepal.Length, and Sepal.Width
- Use logistic regression and experiment with different combinations of the features
- Use 3-fold cross-validation to evaluate models
  - Which combination of features gives the highest precision/recall?

```
from sklearn.linear_model import LogisticRegression
X = iris.data[["petal width (cm), ???"]].values
y = iris.target_names[iris.target] == 'virginica'
```

```
log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)
log_reg.predict(..., ...)
```

# Textbook code

- [Textbook code](#)
- [Textbook code on Google Colab](#)
- Open `04_training_linear_models.ipynb`

# Acknowledgement

- Many slides from Dr. Christopher Ryu
- Content based on “Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow,” Aurélien Géron, 3rd Edition (October 2022), O'Reilly Media, Inc.